

Speaker Notes — Thesis Defense Presentation

HeatsinkLab Wind Tunnel — Automated Heatsink Thermal Characterisation

Alexander Groenvynck

Academic year 2025–2026

Companion to `presentation.pdf` (18 slides). One section per slide. Target: 15–20 minutes total. Approximate timings are noted per slide.

Slide 1 — Title (*~30 s*)

Good morning / afternoon, everyone. My name is Alexander Groenvynck. Today I'll be presenting my bachelor thesis on the HeatsinkLab Wind Tunnel — an automated platform for characterising the thermal performance of heatsinks in a teaching laboratory. I'll walk you through the problem I set out to solve, the system I designed and built, how the software pieces talk to each other, and the results I obtained. The presentation will take around 15 to 20 minutes.

Slide 2 — Table of Contents (*~15 s*)

Here is a quick overview of what I'll cover. I'll start with the motivation and the engineering problem, then walk through the hardware and software design, show the results, and finish with the conclusions and next steps.

Slide 3 — The Problem (*~90 s*)

The core problem is simple: measuring how well a heatsink dissipates heat is fundamentally a waiting game. You heat something up, hold it at a stable temperature, wait for equilibrium, then read the instruments and note everything down — and repeat this for every temperature point and every heatsink. In a teaching lab with 10 or 15 students, that process is slow, inconsistent, and distracts from the actual engineering learning objective, which is comparing heatsinks, not operating instruments. On top of that, you have a live heater running the whole time. So the goal I set was: can I reduce the entire procedure to a single button press, with the heater shutting off automatically when the test is done?

Slide 4 — Research Questions (*~60 s*)

These four research questions guided the whole project. RQ1 is about whether a cheap ESP32 microcontroller and a PID loop are good enough. RQ2 is the automation question — instead of a human watching a thermometer and deciding “OK, it has settled”, can the system do that itself? RQ3 is the usability and safety question together. And RQ4 is the validation question — does the output actually mean something, or are the numbers just noise? I'll answer all four at the end.

Slide 5 — System Overview (Three-Tier Architecture) (~90 s)

The whole system is split into three clear tiers, each talking to the next through a well-defined interface. At the bottom — closest to the hardware — is the firmware running on the ESP32-S3. It runs the PID loop, reads the sensors, and streams telemetry over USB serial at 115,200 baud. The key design decision here was to treat the firmware as a stable black box. I never changed it once it was validated. All user-facing features live in the two software tiers above. The backend is a Python FastAPI server that reads the serial port, parses the telemetry, and relays it to the browser over WebSocket. The frontend is a single HTML file with JavaScript — no build step, no dependencies, just open the file in a browser. I'll go much deeper into how these layers talk to each other in the software slides.

Slide 6 — Wind Tunnel Construction (~90 s)

The wind tunnel is a simple chipboard box. The fan sits at one end, and a hexagonal honeycomb grid sits immediately behind it. The reason for the HEX grid is flow conditioning. A raw axial fan produces a swirling, non-uniform flow — the air spins as it comes off the blades. If you put a heatsink in that, one side sees more airflow than the other, and your measurement is not repeatable. The honeycomb cells break up that rotation and straighten the flow into a more uniform profile across the cross section. This is the same principle used in research and aerodynamic wind tunnels, just at a much smaller scale and lower budget. The result is that every heatsink you test sits in essentially the same airflow conditions — which is the whole point.

Slide 7 — The Heater Module (~90 s)

The heater module is the heart of the physical platform. Working from the bottom: first a layer of insulation to stop heat escaping downward. Then a steel plate for structural rigidity. Then the ceramic PTC heater element — PTC stands for Positive Temperature Coefficient, which means its resistance rises with temperature, giving it a self-limiting property. On top of that is the copper spreader plate, which I want to emphasise: copper has excellent thermal conductivity, and the plate distributes the heat from the relatively small heater element uniformly across the larger base area that the heatsink will sit on. The heatsink clips on top with a spring-steel wire clip, so you can swap heatsinks quickly without tools. The thermocouple tip sits in a drilled hole in the copper — not on the heatsink surface — so the temperature we measure is the copper plate temperature, which is consistent across all tested heatsinks.

Slide 8 — Electronics & Sensors (~90 s)

Let me quickly run through the electronics. The ESP32-S3 is the brain — it reads sensors, runs the PID loop, and talks to the PC. The MAX6675 converts the thermocouple voltage to a digital reading over SPI — it also handles cold-junction compensation internally, which is important for thermocouple accuracy. The INA226 is particularly useful here: it measures both voltage and current simultaneously, so we get actual electrical power dissipated in the heater in watts — not estimated from a setpoint, but measured directly. The two MOSFETs are just electronic switches driven by PWM signals from the ESP32 — one for the heater, one for the fan. Everything sits on a handmade perfboard in a compartment under the tunnel. You can see the rainbow ribbon cable carrying the SPI thermocouple lines, and the red twisted cables for the 12 V heater power.

Slide 9 — Control Strategy: SMART Mode (~90 s)

SMART mode is the key control innovation in the firmware. It starts in PID mode — continuously adjusting heater power to drive the temperature toward the setpoint. Once the temperature stays within half a degree of the setpoint for 8 consecutive seconds, it switches to HOLD: the PWM is frozen at its current value. This is important because a locked PWM means a stable, known power level — which is exactly what we need to compute thermal resistance. The EMA filter smooths the thermocouple reading to reduce cycle-to-cycle noise. There's also a second check on the backend side: the backend tracks the equilibrium PWM estimate and only records a data point when that estimate

has also converged, meaning it changes by less than half a PWM unit between telemetry frames. This dual-check approach means we never record a point prematurely.

Slide 10 — Thermal Metric: R_{th} (~60 s)

The thermal metric is beautifully simple: temperature rise divided by power. If you put 10 watts into a heatsink and its temperature rises 30 degrees above ambient, the thermal resistance is 3 degrees per watt. A better heatsink — one that sheds heat more efficiently — rises less, so its R_{th} is lower. We measure this at several different temperatures to confirm it's stable — R_{th} is a property of the heatsink geometry and material, not of the operating point, so if our measurements are good, the values should be consistent across set-points. The comparison is fair because the wind tunnel runs at the same fan speed for every heatsink, so they all see the same airflow conditions. I want to quickly address the fins vs copper point, because it surprises people: copper has roughly twice the thermal conductivity of aluminium. But under forced airflow, what matters most is convective surface area — how much surface the air can actually touch. A finned aluminium heatsink has many times more surface area than a solid copper block of the same base footprint, so the fins win despite the material disadvantage.

Slide 11 — Software Architecture: Data Flow (~120 s)

This is the full data pipeline from sensor to screen. Let me walk through it left to right. The thermocouple talks to the ESP32 over SPI — that's a fast synchronous bus, the reading takes a few microseconds each firmware cycle. The firmware runs the PID calculation, applies the EMA filter, and emits one telemetry line over USB serial at 115,200 baud — roughly every 100 milliseconds. The backend reads that serial stream continuously, parses each line with a single regular expression called `TELEM_RE`, extracts all the fields, and immediately broadcasts a WebSocket message to every connected browser. The frontend receives that WebSocket message on every cycle and updates the live Chart.js graph and the sidebar readouts. It also runs the tester state machine — I'll explain that shortly. The backend writes every single telemetry frame to the raw log CSV as it arrives. The frontend writes the results CSV and triggers an automatic download when the test completes. The dashed box on the left is the virtual MCU simulator — it produces fake telemetry so you can run the entire workflow without a physical ESP32 attached.

Slide 12 — Backend: `server.py` (~90 s)

The backend is a single Python file — `server.py` — built with FastAPI and the asyncio event loop. The serial reader runs as a continuous async coroutine, reading one text line at a time from the serial port. Each line is handed to `handle_line`, which applies the `TELEM_RE` regular expression — a single, comprehensive regex with named capture groups for every telemetry field. This is deliberately a single regex rather than multiple passes, so there's one place to maintain when the firmware adds a new field. The broadcast function then serialises all the extracted fields into a JSON message and sends it to every connected browser over WebSocket. It also writes the frame to the raw log CSV. The `sendCmd` function queues serial commands — things like `SET SP 50` to set a 50-degree setpoint, or `SET RUN OFF` to cut heater power. The virtual simulator was incredibly valuable: it let me write and debug the entire automated test workflow, the equilibrium detection, the CSV export, and the safety shutdown, all without needing the physical hardware to be ready. Every safety mechanism was tested in simulation before I ever energised the heater.

Slide 13 — Frontend: `index.html` (~90 s)

The frontend is intentionally a single HTML file with everything inline — HTML structure, CSS styles, and all the JavaScript logic. This is a deliberate design choice for the lab deployment: you can copy it to any machine with Python installed and it just works. No package manager, no build step, no internet. The WebSocket message handler receives the telemetry JSON on every firmware cycle and does two things: updates the live Chart.js graph with the new temperature point, and refreshes all the sidebar readout fields. The tester state machine is the key piece of automation. It lives in JavaScript and it watches the incoming telemetry. When it's waiting for equilibrium,

it checks whether the `eq_pwm` field — the equilibrium PWM estimate from the firmware — has converged: changed by less than 0.5 units between frames. When that condition is true, it records the current temperature, power, and other readings as one row in the results table, then either commands the next setpoint or — if that was the last one — triggers the shutdown sequence and automatically downloads the CSV file. The student never has to click Download.

Slide 14 — The Automated Workflow (~90 s)

This slide is really the point of the whole project. Eight steps, but the student only actively does five of them — the rest are automatic. The connection wizard shows a numbered layout so even a student who has never used lab software knows what to do next. Once Start Test is pressed, the system drives the heater to the first setpoint, waits for equilibrium in SMART mode, records the data point, advances to the next setpoint, and repeats until all setpoints are done. Then it shuts off the heater, runs a cool-down, and automatically downloads the CSV. The student doesn't need to watch the screen. They can step away, come back when the test is done, and their results file is already on the desktop. You can see in the screenshots: the port selector on the top left, the live temperature chart in the middle, and the results panel with download buttons at the bottom right.

Slide 15 — Results: R_{th} Comparison (~90 s)

These are representative values that illustrate the format of the output. The first thing to notice is that R_{th} is consistent across set-points for each heatsink — 2.38, 2.42, 2.38 for the fin array. That consistency is exactly what we expect from theory, and it tells us the measurement is working correctly. The second observation is that the finned aluminium heatsink is substantially better than the solid copper block — 2.4 versus 4.0 degrees per watt, a factor of 1.7. Despite copper's higher thermal conductivity, the fins create so much more surface area for the air to contact that they win decisively. Notice also that HS-01 requires more heater power for the same temperature — because it's shedding more heat to the airflow, the heater has to work harder to maintain the setpoint. That's a really clean physical consistency check on the data. The platform answered RQ4: yes, it meaningfully distinguishes heatsinks.

Slide 16 — Safety: Heater Shutdown (~60 s)

Safety was requirement R5 and I treated it as P0 — the highest priority. A live heater element in a teaching lab is a real hazard. I tested three ways the test can end. First, the happy path: test completes normally, backend sends SET RUN OFF, the GUI confirms the heater is off. Second, mid-run stop: pressing the Stop button sends the command immediately, and the INA226 current reading drops to zero within one 100-millisecond firmware cycle. Third, the browser just gets closed or the network drops — this is the failure mode that's easiest to miss. The backend detects the WebSocket disconnection event and sends SET RUN OFF as a fallback within a second or two. In all three cases I confirmed with the INA226 that heater power actually went to zero — not just that the command was sent. After the shutdown, there's a 60-second cool-down period where the fan stays on at low speed, which protects the heater element and the copper plate from thermal shock.

Slide 17 — Limitations & Future Work (~60 s)

Every system has limitations and this one is no different. The most significant technical gap is that airspeed is not measured — we set the fan to a percentage and assume that's reproducible. For comparison between sessions or between setups, you really want a measured air velocity. The good news is that the SDP510 differential pressure sensor is already physically installed and the data schema already has placeholder columns for it — wiring it up and writing the parser update is the next concrete step. The hardware safety watchdog is the other important gap: right now if the ESP32 firmware crashes or hangs, there's nothing to cut the heater power. A small second microcontroller with a watchdog timer and a relay is the correct fix. For future work, the repeatability study is

important for any real scientific use of the platform — we'd need to quantify how much the R_{th} varies between repeated tests of the same heatsink under the same conditions.

Slide 18 — Conclusion (*~60 s*)

To close: I set out to build a system where a student could characterise a heatsink with a single button press, with the heater shutting off automatically. I achieved that. All four research questions have a positive answer. The PID loop is stable enough, equilibrium detection is reliable, the workflow is one button, and the platform clearly distinguishes heatsinks with different thermal properties. The key architectural decision — treating the firmware as a locked black box and putting everything student-facing in the software — meant I could develop, debug, and iterate quickly, and the safety-critical control loop never needed to change. The platform is usable by students who have never heard of PID, don't know what a COM port is, and have no data acquisition background. That was the goal from day one, and the system delivers it. I'm happy to take questions — either on the hardware design, the software architecture, the control strategy, or the results.